

FISTA algorithm: so fast ?

Mathis Scheffler and Constantin Vaillant-Tenzer

January 8, 2025

Objectives of the FISTA Algorithm

- Minimize a function of the form:

$$\min_x \{F(x) = f(x) + g(x)\}$$

- $x \in \mathcal{R}^d$
- $f(x)$: convex and smooth (differentiable) function with L-Lipschitz gradient
- $g(x)$: convex, and lower-semi continuous (potentially non-smooth and non continuous function) function

Algorithm Steps : FISTA

Algorithm 1 FISTA

Initialize: $y_1 = x_0$, $t_1 = 1$

for $k = 1, 2, \dots$ **do**

$$x_k = \text{prox}_{\lambda g}(y_k - \lambda \nabla f(y_k))$$

$$t_{k+1} = \frac{1 + \sqrt{1 + 4t_k^2}}{2}$$

$$y_{k+1} = x_k + \left(\frac{t_k - 1}{t_{k+1}}\right)(x_k - x_{k-1})$$

end for

$$\text{Prox}_{\gamma, g}(\tau) := \arg\min_{\theta \in \Theta} \left(g(\theta) + \frac{1}{2\gamma} \|\theta - \tau\|^2 \right)$$

What happens if ∇f is untracktable ?

- Suppose that ∇f is untracktable and that :

$$\nabla f(\theta) = \int_X H(\theta, x) \pi_\theta(dx)$$

- We can approximate stochastically ∇f !

Stochastic Fista or Perturbed Fista (P-FISTA)

- Classic Fista:

$$x_k = \text{prox}_{\gamma g}(y_k - \gamma \nabla f(y_k))$$

- p-Fista:

$$x_k = \text{prox}_{\gamma g}(y_k - \gamma H_{n+1})$$

with H_{n+1} a stochastic approximation of $\nabla f(y_k)$ obtained with MCMC

Does P-FISTA theoretically works ?

Assumptions :

1
$$\mathcal{L} = \operatorname{argmin}_x (f + g) \neq \emptyset$$

2
$$t_0 = 1; t_n \geq 1; \gamma_n \in \left[0, \frac{1}{L}\right]; \tau_n := \gamma_n t_{n-1}^2 - \gamma_{n+1} t_n (t_n - 1) > 0.$$

3 .

With

$$z_n := \theta_n + t_n (\operatorname{Prox}_{\gamma_{n+1}, g}(\vartheta_n - \gamma_{n+1} \nabla f(\vartheta_n)) - \theta_n)$$

and

$$\eta_{n+1} := H_{n+1} - \nabla f(\vartheta_n)$$

Then

The series $\sum_n \gamma_{n+1} t_n \langle z_n - \theta_*, \eta_{n+1} \rangle$ exists for some $\theta_* \in \mathcal{L}$.

Does P-FISTA theoretically works ?

Results:

- ① The algorithm converges to a minimum of $f+g$
- ② Under conditions such as the following :
 - ▶ $(a \in [0, 1], c = 3 - 2a)$ or $(a \in [1, 2[, c = 2 - a)$ and $b > 1$
 - ▶ $t_n = O(n)$
 - ▶ $\gamma_n = \gamma n^{-a}$
 - ▶ $m_n = m(\ln n)^b n^c = \text{Samples to estimate } H_n$

We obtain convergence at rate n^{a-2} . The optimal rate is Fista's rate ie. $0(n^{-2})$ but taking $a = 0$ leads to a huge computation as it would imply that $n^3 = o(m_n)$.

Numerical Results Implementation

Performance of FISTA:

- Slow in Python (3,5 days on an i7-12700K processor) - with $p = 10$ and $N = 25$.
- Bottleneck: Large sample size.

Processing Time Distribution:

- Algo 1 (P-PG, $m_n = O(\sqrt{n})$, $t_n = 1$): 4%
- Algo 2 (P-FISTA, $m_n = O(n^3)$, $t_n = O(n)$): 19%
- Algo 3 (P-FISTA, $m_n = O(n^3)$, $t_n = O(\sqrt{n})$): 31%
- Algo 4 (P-FISTA, $m_n = O(n^3)$, $t_n = O(n^\epsilon)$): 45%
- Algo 5 (P-PG, $m_n = 10$, S_n computed iteratively): 1%

Limitations:

- Convergence threshold: 10^{-2} (not all 2000 steps completed).
- Sample cap: 10,000 (growth stops after 22 iterations in algorithms 2–4).

Optimization Attempts:

- Parallelization with JAX: Faster but limited.
- Potential C implementation: Estimated $\times 87$ speedup (Heer and al. 2023).

Numerical analysis

Study of the sparsity of the non-zero components of θ_n :








Figure: [left]: number of non-zeros components θ_n , [right]: probability to be non-nul for each component avec n iterations.

Our code is available here: https://github.com/cvt8/fista_sofast.

References I

-  Beck, Amir and Marc Teboulle (2009). “A fast iterative shrinkage-thresholding algorithm for linear inverse problems”. In: *SIAM journal on imaging sciences* 2.1, pp. 183–202.
-  Fort, Gersende et al. (2018). “Stochastic FISTA algorithms: so fast?” In: *2018 IEEE Statistical Signal Processing Workshop (SSP)*. IEEE, pp. 796–800.
-  Heer, Niklas and et al. (2023). *Speed Comparison of Programming Languages*. Accessed: 2025-01-05. URL: <https://github.com/niklas-heer/speed-comparison>.
-  Wolff, Ulli (1989). “Collective Monte Carlo updating for spin systems”. In: *Physical Review Letters* 62.4, p. 361.

Questions



Metropolis-Hastings Algorithm

- Goal: Sample from a target distribution $p(x)$ when direct sampling is difficult
- Algorithm Steps:
 - 1 Start with initial state x_t
 - 2 Propose new state x' from proposal distribution $q(x'|x_t)$
 - 3 Calculate acceptance ratio:

$$\alpha = \min\left(1, \frac{p(x')q(x_t|x')}{p(x_t)q(x'|x_t)}\right)$$

- 4 Accept x' with probability α :
 - If accepted: $x_{t+1} = x'$
 - If rejected: $x_{t+1} = x_t$
- Key feature: Only needs to know $p(x)$ up to a normalizing constant
 - Converges to the target distribution as $t \rightarrow \infty$

Gibbs Sampling

- Special case of MH with acceptance rate always 1
- Used when sampling from conditional distributions is easier
- For variables $\mathbf{x} = (x_1, \dots, x_n)$:
 - 1 Initialize $\mathbf{x}^{(0)}$
 - 2 For each iteration t :

$$x_1^{(t+1)} \sim p(x_1 | x_2^{(t)}, x_3^{(t)}, \dots, x_n^{(t)})$$

$$x_2^{(t+1)} \sim p(x_2 | x_1^{(t+1)}, x_3^{(t)}, \dots, x_n^{(t)})$$

$$\vdots$$

$$x_n^{(t+1)} \sim p(x_n | x_1^{(t+1)}, \dots, x_{n-1}^{(t+1)})$$

- Advantages:
 - ▶ No tuning of proposal distributions needed
 - ▶ Higher acceptance rate than MH
 - ▶ Works well for high-dimensional problems

Approximate Bayesian Computation (ABC)

Key Idea: Approximate posterior when likelihood is intractable

Algorithm:


- 1 Sample $\theta^* \sim p(\theta)$ from prior
- 2 Simulate data $D^* \sim p(D|\theta^*)$
- 3 Compute summary statistics $S^* = S(D^*)$
- 4 Accept if $\rho(S^*, S_{obs}) \leq \varepsilon$

Key Components:

- Summary statistics $S(\cdot)$
- Distance metric $\rho(\cdot, \cdot)$
- Tolerance threshold ε

Variants:

- ABC-MCMC: Uses Markov Chain exploration



Wolf Sampling (Wolff 1989)

Context: Designed for Monte Carlo simulations of spin-lattice systems (e.g., Ising model).

Uses cluster updates to improve efficiency.

Main Steps:

- ➊ **Initialization:** Select a random spin s_0 as the cluster seed.
- ➋ **Cluster Growth:**
 - ▶ Add aligned neighboring spins to the cluster with probability $P_{\text{add}} = 1 - e^{-2\beta J}$, where β is the inverse temperature and J the interaction energy.
- ➌ **Cluster Flip:** Flip all spins in the cluster.

Advantages:

- Reduces correlations between successive configurations.
- Effective near the critical temperature.

Numerical Simulations: Overview

- Goal: Compare five algorithms in optimizing a likelihood function in binary graphical models.
- Setting:
 - ▶ Sparse matrix θ with $p = 100, N = 250$.
 - ▶ Penalty term: $g(\theta) = \lambda \sum |\theta_{ij}| + \mu \sum \theta_i^2$.
- Algorithms:
 - ▶ Alg1: Perturbed Proximal Gradient (P-PG).
 - ▶ Alg2–4: Variants of Perturbed FISTA (P-FISTA).
 - ▶ Alg5: Averaging-based Perturbed Proximal Gradient.

Algorithm 1: Perturbed Proximal Gradient (P-PG)

- Update rule:

$$\theta_{n+1} = \text{Prox}_{\gamma_n g}(\theta_n - \gamma_n H_n),$$

where H_n is a Monte Carlo approximation of the gradient.

- Parameters:

- ▶ Step size: $\gamma_n = O(1/\sqrt{n})$.
- ▶ Sample size: $m_n = O(\sqrt{n})$.

- Results:

- ▶ Convergence rate: $O(1/n)$ (logarithmic terms ignored).
- ▶ Baseline for comparison.

Algorithm 2: P-FISTA with $t_n = O(n)$

- Accelerated update rule:

$$\theta_{n+1} = \text{Prox}_{\gamma_n g}(\varphi_n - \gamma_n H_n),$$

where $\varphi_n = \theta_n + \frac{t_{n-1}-1}{t_n}(\theta_n - \theta_{n-1})$.

- Parameters:

- ▶ Step size: $\gamma_n = O(1)$.
- ▶ Sample size: $m_n = O(n^3)$.
- ▶ Momentum: $t_n = O(n)$.

- Results:

- ▶ In practice, slower convergence than Alg1.
- ▶ Optimal convergence rate $O(1/n^2)$.

Algorithm 3: P-FISTA with $t_n = O(\sqrt{n})$

- Same update rule as Algorithm 2.
- Parameters:
 - ▶ Step size: $\gamma_n = O(1)$.
 - ▶ Sample size: $m_n = O(n^3)$.
 - ▶ Momentum: $t_n = O(\sqrt{n})$.
- Results:
 - ▶ Convergence slower than $t_n = O(n)$ but better than Alg1.
 - ▶ Computational cost remains high due to $m_n = O(n^3)$.

Algorithm 4: P-FISTA with $t_n = O(n^\epsilon)$

- Same update rule as Algorithm 2.
- Parameters:
 - ▶ Step size: $\gamma_n = O(1)$.
 - ▶ Sample size: $m_n = O(n^3)$.
 - ▶ Momentum: $t_n = O(n^\epsilon), \epsilon \ll 1$.
- Results:
 - ▶ Practical trade-off between convergence speed and stability.
 - ▶ The slower convergence in practice
 - ▶ Requires further theoretical analysis.

Algorithm 5: Averaging-based P-PG

- Update rule:

$$S_{n+1} = (1 - \delta_{n+1})S_n + \delta_{n+1} \frac{1}{m_n} \sum_{j=1}^{m_n} S(X_j),$$

where $\delta_n = O(n^{-0.9})$.

- Parameters:

- ▶ Sample size: $m_n = O(1)$.
- ▶ Uses all past samples iteratively.

- Results:

- ▶ Improves convergence rate significantly.
- ▶ Best performance in the simulations.